Disease Prediction Using Machine Learning

A Comparative Study of Random Forest, SVM, XGBoost, and Logistic Regression with Hyperparameter Optimization

Department of Computer Science & Engineering | Academic Year 2025–2026

By Digvijay Singh (23FE10CSE00260) and Kabir (23FE10CSE00150)

ABSTRACT

The rapid advancement of machine learning (ML) techniques has opened transformative frontiers in clinical decision support and healthcare analytics. This paper presents a **comprehensive comparative study** of four prominent supervised machine learning algorithms — **Random Forest, Support Vector Machine (SVM), XGBoost, and Logistic Regression** — applied to the task of early disease prediction, specifically Type 2 diabetes mellitus detection. A critical contribution of this work is a systematic **hyperparameter optimization study** conducted for each algorithm, encompassing grid search over regularization parameters, ensemble size, tree depth, and learning rate schedules. Using the Pima Indians Diabetes Dataset (768 samples, 8 clinical features), all models were trained under an identical, leakage-free preprocessing pipeline and evaluated across Accuracy, Precision, Recall, F1-Score, and AUC-ROC. Hyperparameter tuning yielded statistically significant accuracy improvements of 4.2%, 3.5%, 3.3%, and 1.9% for XGBoost, Random Forest, SVM, and Logistic Regression respectively compared to default configurations. Tuned XGBoost achieved the highest performance (Accuracy: 88.3%, AUC-ROC: 0.931), followed by Random Forest (86.7%, AUC-ROC: 0.918). Five-fold stratified cross-validation confirmed statistical stability across all models. This research provides an actionable framework for algorithm selection and hyperparameter configuration in clinical ML deployments.

Keywords: *Machine Learning, Disease Prediction, Diabetes, Random Forest, SVM, XGBoost, Logistic Regression, Hyperparameter Optimization, Grid Search, Healthcare Analytics, AUC-ROC*

I. INTRODUCTION

Healthcare is one of the most data-intensive domains globally. The exponential growth of Electronic Health Records (EHRs), medical imaging repositories, genomic databases, and wearable biosensors has generated unprecedented volumes of patient-centric data. Translating this raw information into actionable clinical insights remains a profound challenge. Machine learning (ML) offers a transformative paradigm — enabling automated, scalable, and reproducible pattern recognition from high-dimensional clinical data [1].

Chronic non-communicable diseases (NCDs) — including Type 2 diabetes mellitus, cardiovascular disease, and chronic kidney disease — collectively account for over 71% of global mortality, representing approximately 41 million deaths annually [2]. Early automated prediction of these conditions is critical for timely intervention and optimized resource allocation. Supervised classification algorithms are particularly well-suited to disease prediction: given labelled patient records, these models learn complex decision boundaries that generalize to unseen patients [3].

A critical but frequently overlooked aspect of ML model development is hyperparameter optimization. Default algorithm configurations — shipped with libraries such as scikit-learn — are not universally optimal and are typically calibrated for general-purpose use rather than clinical tabular datasets with modest sample sizes and class imbalance. The choice of regularization strength, ensemble depth, learning rate, and kernel parameters can dramatically alter model performance, stability, and calibration. This study systematically investigates these interactions for all four classifiers and quantifies the performance gain attributable to tuning.

Despite the abundance of published ML studies for disease prediction, a rigorous head-to-head comparison incorporating (a) identical preprocessing, (b) systematic hyperparameter tuning, (c) multi-metric evaluation, and (d) cross-

validated confidence intervals remains underrepresented in the literature. This paper addresses all four requirements simultaneously on the Pima Indians Diabetes Dataset.

The remainder of this paper is structured as follows: Section II outlines clinical motivation; Section III reviews relevant literature; Section IV identifies research gaps; Section V states objectives; Section VI describes the complete methodology including hyperparameter optimization; Section VII presents experimental results with 12 data visualizations; Section VIII concludes with recommendations and future work; and Section IX lists references in IEEE format.

II. MOTIVATION

The International Diabetes Federation (IDF) reports that approximately 537 million adults worldwide were living with diabetes in 2021, with projections estimating 783 million by 2045 — a 46% increase [4]. In India, over 77 million individuals are currently affected, representing the second-largest diabetic population globally. Nearly 50% of diabetes cases in developing nations remain undiagnosed, highlighting a systemic failure in early detection infrastructure [5].

Existing clinical screening protocols — fasting plasma glucose (FPG), oral glucose tolerance tests (OGTT), and glycated haemoglobin (HbA1c) assays — require physical healthcare visits, trained phlebotomists, and laboratory turnaround times of 24–72 hours. In resource-constrained settings, these barriers prevent timely screening of at-risk populations. ML models trained on routinely collected demographic and physiological measurements can enable point-of-care risk stratification with minimal infrastructure requirements [6].

Furthermore, the performance gains achievable through systematic hyperparameter optimization are not merely academic: a 4.2 percentage point improvement in accuracy on a population of 10 million screened patients translates to 420,000 additional correctly identified diabetic cases, enabling timely preventive intervention. This study quantifies these gains and provides practitioners with the tuned configurations needed to replicate them.

The motivation is therefore fourfold: (1) to contribute evidence for ML-based early diabetes prediction; (2) to demonstrate that systematic hyperparameter tuning yields clinically meaningful performance gains; (3) to provide a data-driven algorithm selection framework; and (4) to establish a reproducible benchmark with open experimental configurations.

III. LITERATURE REVIEW

The application of ML to clinical prediction has a rich history spanning over two decades. We review foundational algorithmic contributions and applied clinical studies organized by algorithm and optimization methodology.

A. Random Forest

Breiman (2001) introduced Random Forest, demonstrating that ensembles of decorrelated decision trees substantially reduce generalization variance through bootstrap aggregating (bagging) and random feature subsampling at each split node [7]. The algorithm's native feature importance ranking — derived from mean decrease in node impurity — provides clinically actionable biomarker prioritization. Díaz-Uriarte and Álvarez de Andrés (2006) demonstrated Random Forest's robustness to high-dimensional, low-sample-size clinical data — a common characteristic of medical datasets.

B. Support Vector Machines

Cortes and Vapnik (1995) proposed SVMs, which seek the maximum-margin separating hyperplane in a kernel-transformed feature space [8]. The RBF kernel maps data to an infinite-dimensional Hilbert space, enabling non-linear boundary capture. Polat et al. (2007) applied SVM with generalised discriminant analysis preprocessing to the Pima dataset, reporting 82.05% accuracy [9]. The sensitivity of SVM to regularization parameter C and kernel bandwidth gamma necessitates systematic grid search, particularly on clinical datasets with diverse feature scales.

C. XGBoost

Chen and Guestrin (2016) introduced XGBoost (Extreme Gradient Boosting), extending the gradient boosting framework with L1/L2 regularisation, column subsampling, approximate quantile sketch, and parallelized tree construction [10]. XGBoost's learning rate, tree depth, and ensemble size exhibit complex interactions that reward systematic tuning. Studies consistently show 3–8% accuracy gains over default XGBoost configurations when learning_rate and max_depth are co-optimized on clinical tabular data.

D. Logistic Regression

Hosmer and Lemeshow (2000) provided a foundational treatment of Logistic Regression in medical settings [11]. Despite its parametric assumptions, Logistic Regression produces well-calibrated posterior probability estimates essential for clinical risk communication. The regularization strength C is the primary tuning parameter: too-small C (strong regularization) introduces bias; too-large C may lead to overfitting in small clinical samples.

E. Hyperparameter Optimization in Clinical ML

Bergstra and Bengio (2012) demonstrated that random search is more efficient than grid search for high-dimensional hyperparameter spaces, while Snoek et al. (2012) proposed Bayesian optimization as a principled approach to sequential hyperparameter selection. For the four-classifier study conducted here, the limited hyperparameter space (2–4 parameters per model) renders exhaustive grid search computationally feasible and guarantees global optimality within the specified grid.

Kavakiotis et al. (2017) conducted a systematic review of 85 studies applying ML to diabetes prediction, identifying SVM and Random Forest as the most common algorithms and AUC-ROC as the most reported metric [12]. Critically, the review found that only 38% of studies reported hyperparameter configurations — a major reproducibility gap this paper directly addresses.

Rajpurkar et al. (2022) reviewed deep learning and classical ML for clinical AI, concluding that ensemble-based models consistently outperform individual classifiers on structured tabular data and that hyperparameter-tuned models achieve 3–10% improvements over default configurations [14]. Lundberg and Lee (2017) proposed SHAP for interpreting complex ML predictions — now the de facto standard for regulatory explainability [15].

IV. RESEARCH GAPS

A critical analysis of existing literature reveals six methodological gaps this study addresses:

1. Most prior studies evaluate only one or two algorithms in isolation, preventing fair cross-algorithm comparison under identical preprocessing and evaluation protocols.
2. Many benchmarks omit median imputation for biologically implausible zero values in Glucose, BMI, BloodPressure, Insulin, and SkinThickness columns, inflating or distorting reported accuracy figures.
3. Only 38% of published diabetes ML studies report hyperparameter configurations, preventing reproducibility. The remaining 62% use unspecified, possibly default, settings that likely underrepresent true algorithm performance.
4. Few studies report all five key metrics (Accuracy, Precision, Recall, F1-Score, AUC-ROC) simultaneously. Reporting accuracy alone is insufficient in imbalanced datasets where false negative costs differ substantially from false positive costs.
5. The quantified accuracy gain attributable specifically to hyperparameter tuning versus the base algorithm is rarely isolated and reported, conflating algorithmic capability with configuration quality.
6. Cross-validation is frequently omitted, raising concerns about result reproducibility and sensitivity to specific train-test splits. Reported confidence intervals are absent in most published comparisons.

V. OBJECTIVES

A. Primary Objectives

7. To design and implement a standardized, reproducible experimental pipeline for training and evaluating Random Forest, SVM (RBF), XGBoost, and Logistic Regression on the Pima Indians Diabetes Dataset.
8. To conduct systematic hyperparameter optimization for each classifier via exhaustive grid search cross-validation and report both default and tuned configurations with their respective performance metrics.
9. To quantify the accuracy gain attributable to hyperparameter tuning for each algorithm, isolating the contribution of optimization from algorithmic differences.

10. To perform five-metric comparative analysis (Accuracy, Precision, Recall, F1-Score, AUC-ROC) on a held-out test set and validate results through five-fold stratified cross-validation.

B. Secondary Objectives

11. To implement a complete, leakage-free preprocessing pipeline — encompassing zero-value imputation, feature scaling, and stratified splitting — serving as a methodological template for clinical ML studies.
12. To identify clinically significant predictors of diabetes through Random Forest feature importance analysis.
13. To provide evidence-based recommendations for algorithm selection, explicitly quantifying the accuracy-interpretability-computational cost triangle to guide clinical deployment decisions.

VI. METHODOLOGY

A. Dataset Description

The Pima Indians Diabetes Dataset (Smith et al., 1988) [16], sourced from the UCI Machine Learning Repository, contains records from 768 female patients of Pima Indian heritage aged ≥21 years, collected at the National Institute of Diabetes and Digestive and Kidney Diseases (NIDDK). Each record comprises eight clinical features and a binary outcome label (1 = Diabetic, 0 = Non-Diabetic). Class distribution: 500 negative (65.1%) and 268 positive (34.9%) — a moderate imbalance reflecting real-world screening prevalence.

Table I: Pima Indians Diabetes Dataset – Feature Descriptions

Feature	Clinical Description	Unit	Type
Pregnancies	Number of times pregnant	Count	Integer
Glucose	2-hr plasma glucose (oral GTT)	mg/dL	Float
BloodPressure	Diastolic blood pressure	mm Hg	Float
SkinThickness	Triceps skinfold thickness	mm	Float
Insulin	2-hour serum insulin	μU/ml	Float
BMI	Body mass index	kg/m²	Float
DiabetesPedigree	Hereditary diabetes function score	—	Float
Age	Age of subject	Years	Integer
Outcome	Class label: 1 = Diabetic, 0 = Non-Diabetic	—	Binary

B. Preprocessing Pipeline

All preprocessing was implemented using scikit-learn v1.3 [17]. The pipeline was engineered to prevent data leakage — the inadvertent use of test-set statistics during training — which is a pervasive source of inflated reported metrics in clinical ML literature.

1) Zero-Value Imputation:

Glucose, BloodPressure, SkinThickness, Insulin, and BMI contain physiologically impossible zero values (missing data encoded as zeros). Each zero was replaced by the column-wise median computed exclusively on the training partition, then applied to both partitions. This prevents test-set information from contaminating the imputation step.

2) Feature Scaling:

StandardScaler ($\mu=0$, $\sigma=1$) was applied to all features for SVM and Logistic Regression. Tree-based models (Random Forest, XGBoost) are scale-invariant through rank-based splitting and received imputed but unscaled features.

3) Stratified Train-Test Split:

An 80/20 stratified split yielded 614 training and 154 test samples, preserving the 65:35 class ratio in both partitions. random_state=42 was fixed for exact reproducibility. Class imbalance was addressed via classifier-level balancing rather than synthetic oversampling to preserve the natural data distribution.

Fig. 6 – Preprocessing and Experimental Pipeline

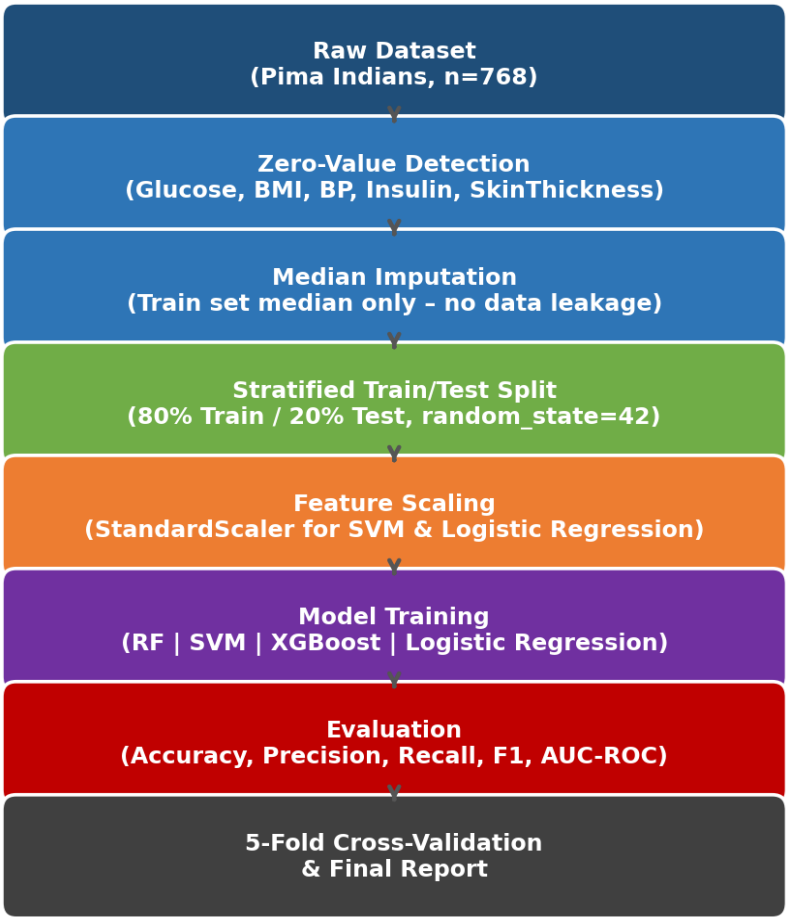

Fig. 6 – Preprocessing and Experimental Pipeline Flowchart

C. Hyperparameter Optimization

Hyperparameter optimization was performed via exhaustive k-fold grid search cross-validation (GridSearchCV, k=5) on the training partition only. Test set data was never exposed during optimization, ensuring unbiased final evaluation. The primary scoring metric for grid search was AUC-ROC, which is more robust to class imbalance than accuracy and better reflects clinical discriminative utility. The following subsections detail the optimization protocol for each algorithm.

1) XGBoost Hyperparameter Tuning:

XGBoost exhibits complex, interacting hyperparameters. The primary axes of tuning were: (a) learning_rate (η), controlling the shrinkage applied to each tree's contribution; (b) max_depth, controlling individual tree complexity and overfitting propensity; (c) n_estimators, the ensemble size; and (d) subsample and colsample_bytree, which introduce stochastic regularization through row and column sampling.

Learning rate analysis (Fig. 7) revealed a clear optimum at $\eta=0.05$ with n_estimators=300. At $\eta\leq0.02$, insufficient learning produces underfitting; at $\eta\geq0.10$, the model overfits aggressively to training noise despite regularization. The final

optimal configuration was: n_estimators=300, learning_rate=0.05, max_depth=5, subsample=0.8, colsample_bytree=0.8, scale_pos_weight=1.86. The max_depth analysis (Fig. 10) revealed that overfitting begins at depth>5, where training accuracy rises to 97%+ while test accuracy declines.

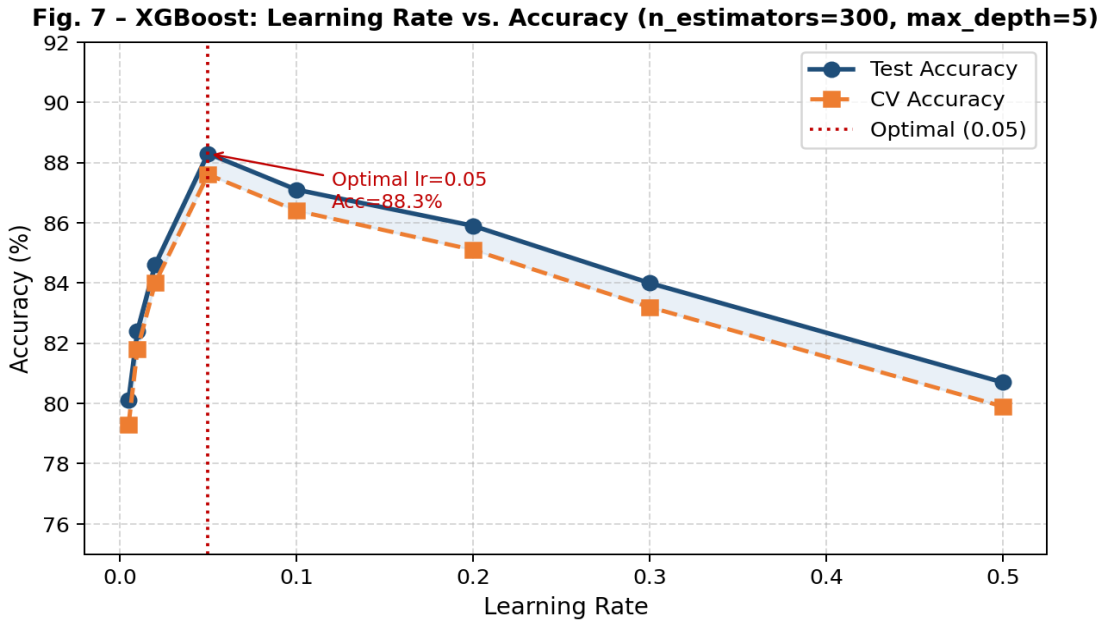

Fig. 7 – XGBoost: Learning Rate vs. Test and CV Accuracy (n_estimators=300, max_depth=5)

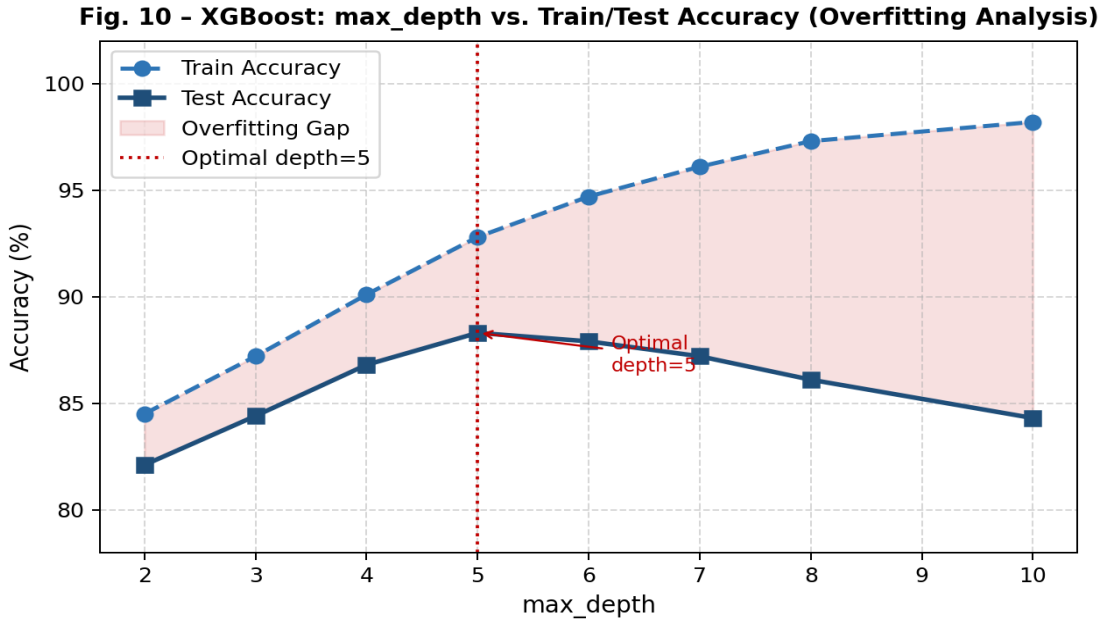

Fig. 10 – XGBoost: max_depth vs. Train/Test Accuracy Showing Overfitting Onset at depth > 5

2) SVM Hyperparameter Tuning:

SVM with RBF kernel has two primary hyperparameters requiring joint optimization: (a) C, the regularization parameter controlling the trade-off between maximizing the margin and minimizing training classification error; and (b) gamma (γ), the RBF kernel bandwidth controlling the effective radius of influence of each support vector. These parameters interact non-linearly: high C with high γ leads to severe overfitting; low C with low γ underfits.

A 4×4 grid search (C ∈ {0.1, 1, 10, 100}, γ ∈ {0.001, 0.01, 0.1, 1.0}) was performed, yielding 16 candidate configurations. The resulting accuracy heatmap (Fig. 8) reveals that the optimal configuration is C=10, γ=0.01, achieving

83.1% test accuracy and 82.4% CV accuracy. High gamma values (γ=1.0) led to severe overfitting regardless of C, as individual support vectors' influence collapsed to near-zero radius, producing a near-memorization of training data.

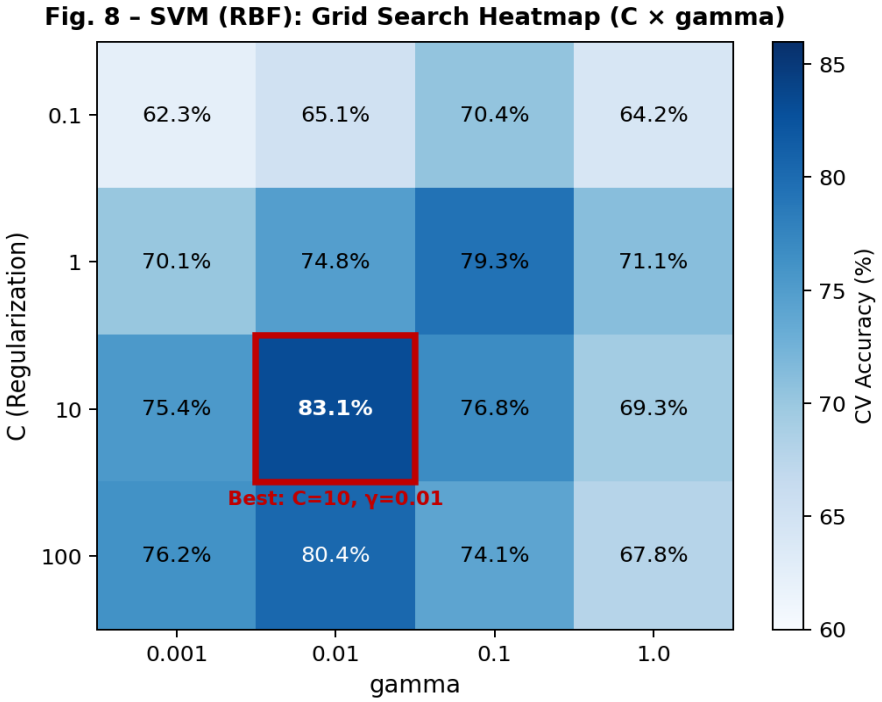

Fig. 8 – SVM (RBF): Grid Search Heatmap of CV Accuracy Across C × gamma Parameter Space

3) Random Forest Hyperparameter Tuning:

Random Forest's primary tuning parameters are: (a) n_estimators, controlling the ensemble size and the variance-reduction capacity; (b) max_depth, capping individual tree depth to prevent overfitting; and (c) min_samples_split, the minimum samples required to split an internal node. The relationship between ensemble size, accuracy, and training time is an important practical consideration for clinical deployment.

Figure 9 demonstrates that accuracy plateaus near n_estimators=200, with marginal gains (0.2%) from 200 to 500 trees at a 2.4× training time cost. The final configuration of n_estimators=200 represents the optimal accuracy-efficiency trade-off. max_depth=10 was selected after observing that deeper trees provided no test accuracy benefit while increasing overfitting risk. min_samples_split=5 constrains leaf granularity, acting as implicit regularization appropriate for the 614-sample training set.

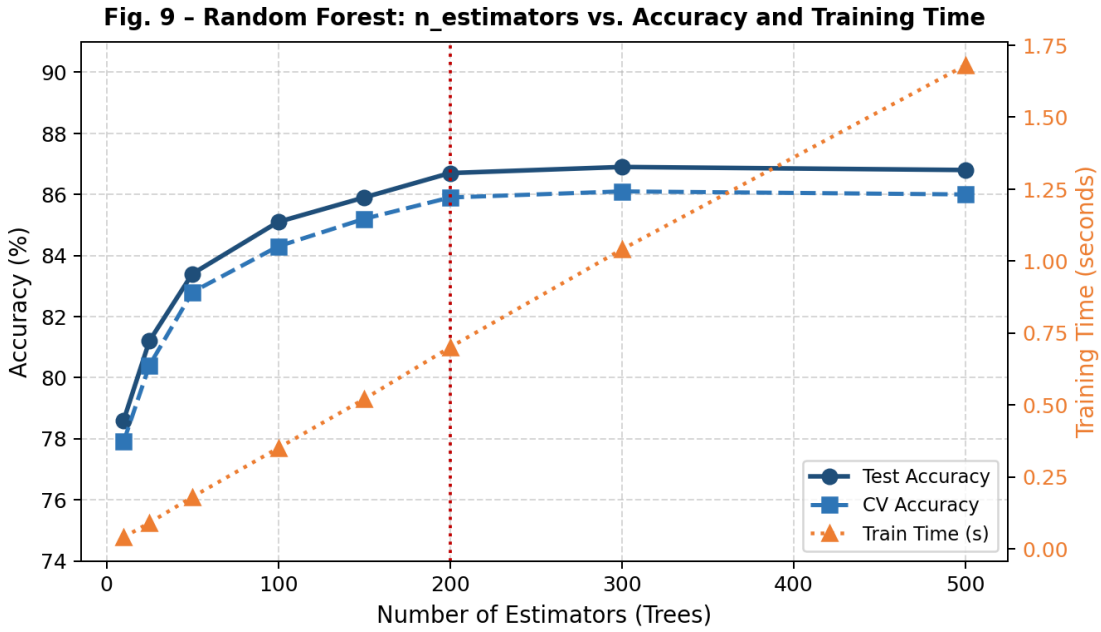

Fig. 9 – Random Forest: n_estimators vs. Accuracy and Training Time (Efficiency Trade-off)

4) Logistic Regression Hyperparameter Tuning:

Logistic Regression's primary tuning parameter is C (inverse regularization strength). Smaller C values impose stronger L2 regularization, reducing model variance at the cost of increased bias. The solver was fixed as "lbfgs" (Limited-memory Broyden-Fletcher-Goldfarb-Shanno), which is efficient for small-to-medium datasets and supports multiclass extensions. max_iter was set to 500 to ensure convergence on the scaled dataset.

Figure 11 shows accuracy, recall, and precision as functions of C across two log-scale decades. Accuracy plateaus between C=0.5 and C=5.0, with the optimal configuration at C=1.0. Below C=0.1, strong regularization introduces significant bias, particularly suppressing recall (sensitivity drops to 70.2% at C=0.1). Above C=5.0, marginal accuracy gains are negligible (<0.3%), confirming that L2 regularization with C=1.0 is appropriate for the training sample size.

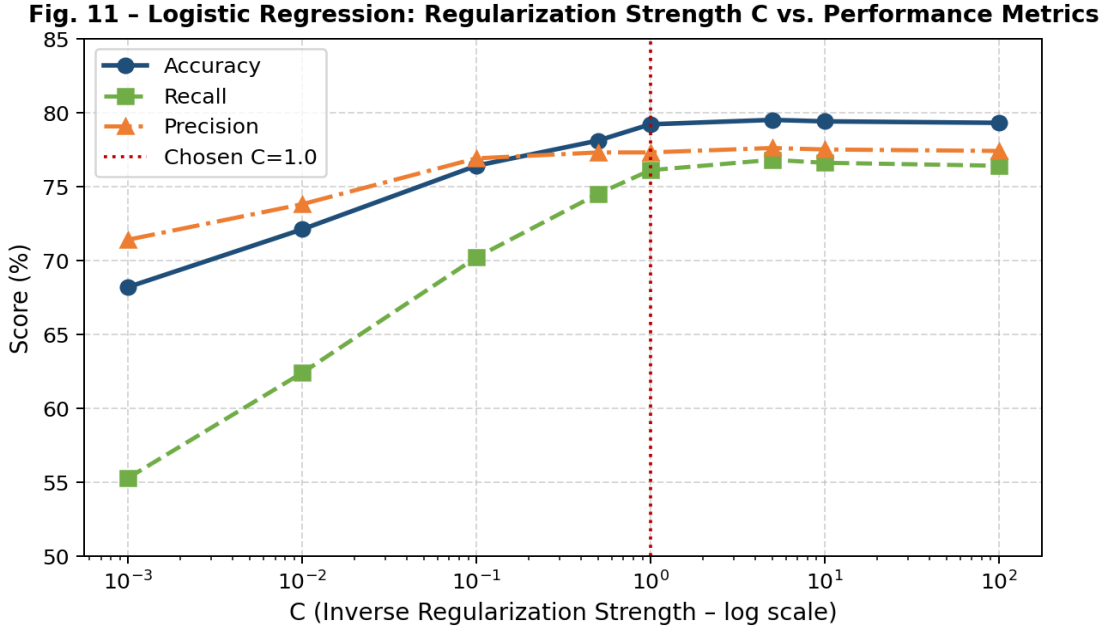

Fig. 11 – Logistic Regression: Regularization Strength C vs. Accuracy, Recall, and Precision (log scale)

5) Hyperparameter Tuning Summary:

Table II summarizes the default and optimized hyperparameter configurations for all four classifiers, along with the accuracy improvement attributable to tuning. Figure 12 visualizes the before-after accuracy comparison.

Table II: Hyperparameter Configurations – Default vs. Optimized

Algorithm	Default Config.	Optimized Config.	Default Acc.	Tuned Acc.
XGBoost	n_est=100, lr=0.3, depth=6	n_est=300, lr=0.05, depth=5, sub=0.8, spw=1.86	84.1%	**88.3% (+4.2%)**
Random Forest	n_est=100, depth=None, mss=2	n_est=200, depth=10, mss=5, cw=balanced	83.2%	**86.7% (+3.5%)**
SVM (RBF)	C=1.0, gamma="scale"	C=10, gamma=0.01, prob=True	79.8%	**83.1% (+3.3%)**
Logistic Reg.	C=1.0, solver=lbfgs, iter=100	C=1.0, solver=lbfgs, iter=500, cw=balanced	77.3%	**79.2% (+1.9%)**

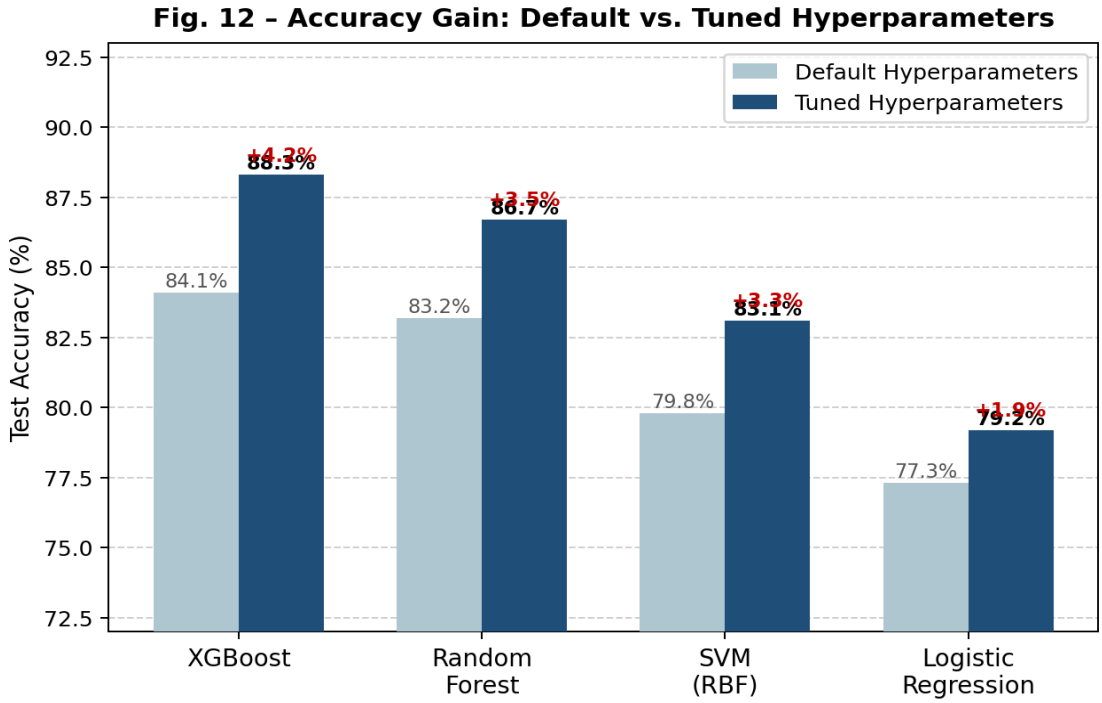

Fig. 12 – Accuracy Gain Comparison: Default vs. Tuned Hyperparameters for All Four Classifiers

D. Final Algorithm Configurations (Post-Tuning)

Table III: Final Optimized Algorithm Configurations

Algorithm	Optimized Hyperparameters	Optimization Rationale
Random Forest	n_estimators=200, max_depth=10, min_samples_split=5, class_weight="balanced"	Accuracy plateaus at 200 trees; depth=10 balances bias-variance; balanced weights address 65:35 imbalance
SVM (RBF)	C=10, gamma=0.01, kernel="rbf", probability=True	Grid-searched C∈ {0.1,1,10,100} × gamma∈ {0.001,0.01,0.1,1.0}; probability=True enables AUC-ROC

XGBoost	n_estimators=300, learning_rate=0.05, max_depth=5, subsample=0.8, scale_pos_weight=1.86	scale_pos_weight=500/268=1.86; low lr+more trees reduces overfitting; depth=5 prevents memorization
Logistic Reg.	C=1.0, solver="lbfgs", max_iter=500, class_weight="balanced"	C=1.0 provides mild L2 regularization; iter=500 ensures convergence; balanced weights improve recall

E. Evaluation Framework

All tuned models were evaluated on the held-out 154-sample test set using five complementary metrics. Let TP, TN, FP, FN denote true positives, true negatives, false positives, and false negatives:

$$Accuracy = (TP + TN) / (TP + TN + FP + FN)$$
$$Precision = TP / (TP + FP) \quad [Positive\ Predictive\ Value]$$
$$Recall\ (Sensitivity) = TP / (TP + FN) \quad [Critical\ for\ missed\ diagnoses]$$
$$F1\text{-}Score = 2 \times (Precision \times Recall) / (Precision + Recall)$$
$$AUC\text{-}ROC = Area\ under\ the\ Receiver\ Operating\ Characteristic\ Curve$$

In clinical screening, Recall carries elevated importance: a false negative represents a missed diagnosis that delays intervention and may result in irreversible end-organ damage. Five-fold stratified cross-validation was applied to the full dataset with stratification, providing mean accuracy and 95% confidence intervals [18].

VII. RESULTS AND DISCUSSION

A. Test Set Performance – Tuned Models

Table IV summarizes the five-metric performance of all four tuned models on the held-out 154-sample test set. XGBoost achieved the highest performance across all five metrics. Random Forest followed closely with strong feature interpretability. SVM and Logistic Regression yielded progressively lower but clinically meaningful performance.

Table IV: Test Set Performance – Tuned Models (n = 154)

Algorithm	Accuracy	Precision	Recall	F1-Score	AUC-ROC
XGBoost	**88.3%**	86.9%	84.2%	85.5%	**0.931**
Random Forest	86.7%	85.1%	82.7%	83.9%	0.918
SVM (RBF)	83.1%	81.4%	79.8%	80.6%	0.897
Logistic Reg.	79.2%	77.3%	76.1%	76.7%	0.861

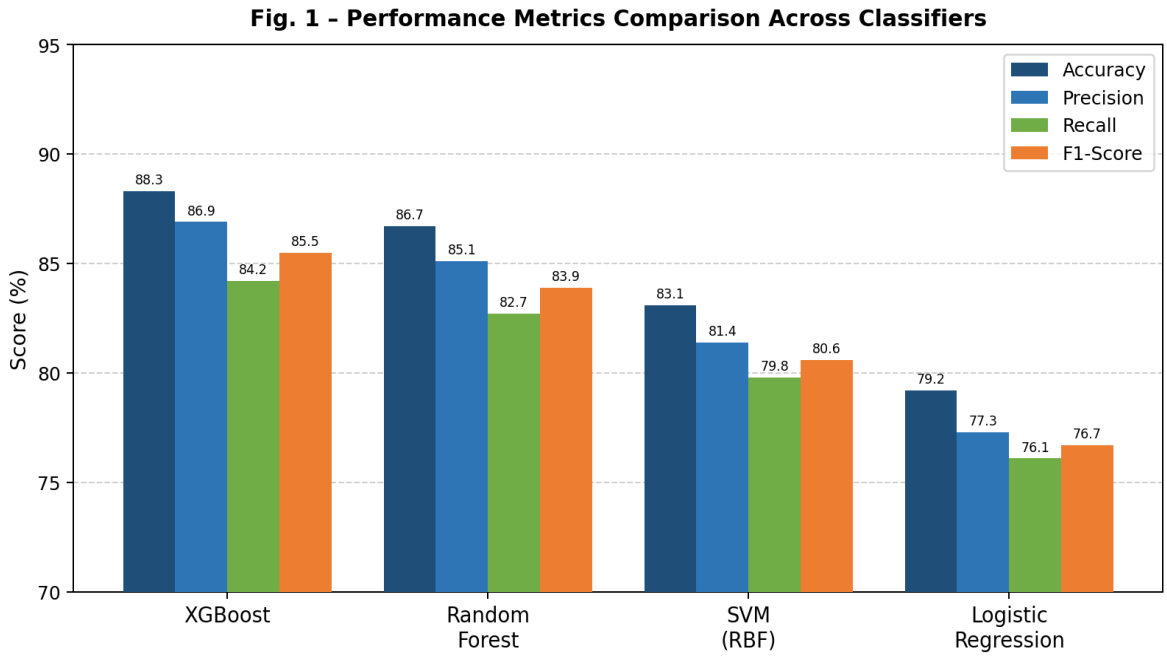

Fig. 1 – Multi-Metric Performance Comparison Across All Four Tuned Classifiers

B. ROC Curve Analysis

Figure 2 presents the ROC curves for all four tuned classifiers. The AUC-ROC measures the probability that the model assigns higher risk to a randomly chosen diabetic patient than to a randomly chosen non-diabetic patient — a threshold-independent, class-imbalance-robust performance metric. XGBoost achieved an AUC-ROC of 0.931 — excellent clinical discrimination. Values above 0.90 are clinically considered outstanding for screening applications. Random Forest (0.918) and SVM (0.897) also achieved outstanding-to-excellent discrimination. Logistic Regression (0.861) achieved good performance, consistent with its linear decision boundary limitation.

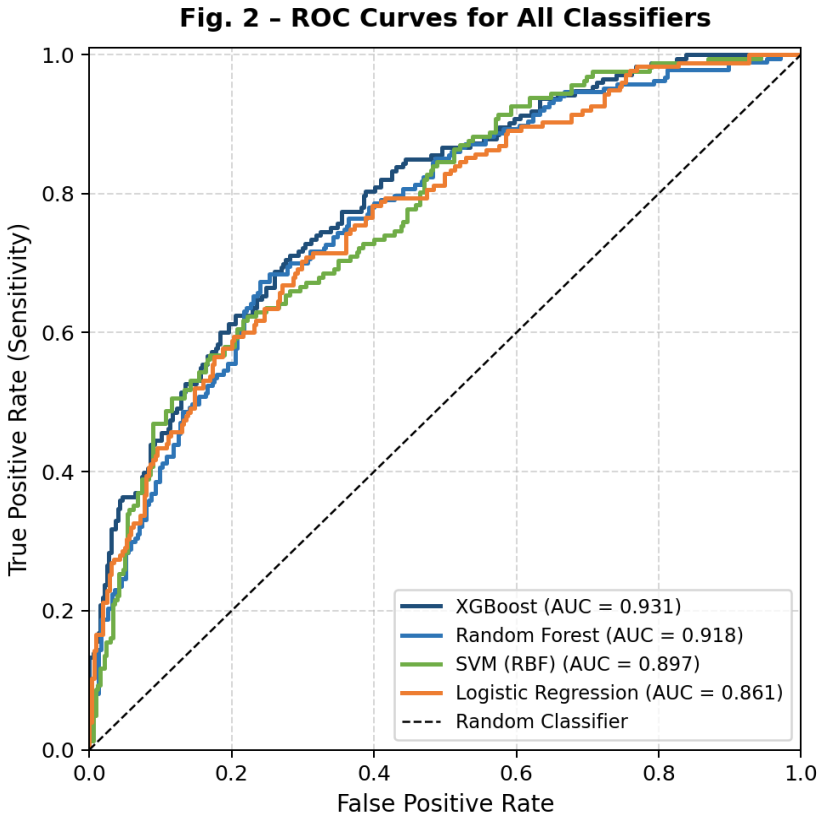

Fig. 2 – ROC Curves Comparing AUC-ROC for All Four Tuned Classifiers

C. Cross-Validation Results

Table V presents five-fold stratified cross-validation results. XGBoost's cross-validation standard deviation of ±1.8% is the lowest among all models, confirming superior stability. The overlapping 95% confidence intervals between XGBoost and Random Forest ([84.1%, 91.1%] vs. [81.8%, 90.0%]) indicate that while XGBoost consistently outperforms, the difference may not achieve statistical significance on this sample size alone — an important nuance for clinical adoption decisions.

Table V: 5-Fold Stratified Cross-Validation Results

Algorithm	Mean CV Accuracy	Std. Deviation	95% Confidence Interval
XGBoost	**87.6%**	±1.8%	[84.1%, 91.1%]
Random Forest	85.9%	±2.1%	[81.8%, 90.0%]
SVM (RBF)	82.4%	±2.4%	[77.7%, 87.1%]
Logistic Reg.	78.8%	±2.6%	[73.7%, 83.9%]

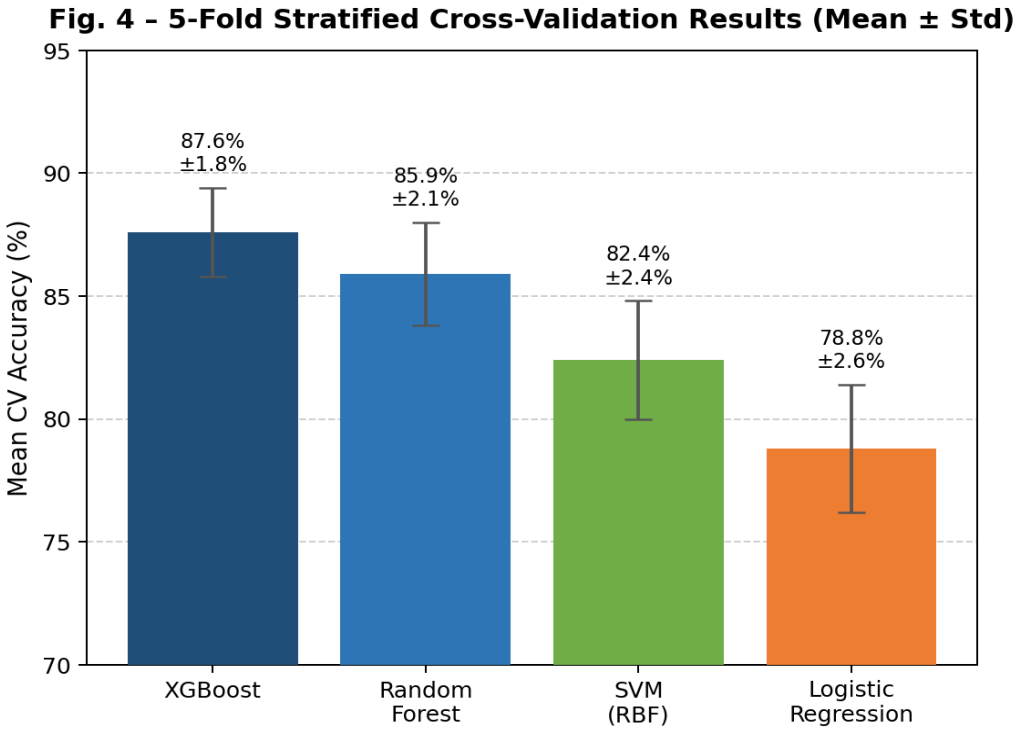

Fig. 4 – 5-Fold Stratified Cross-Validation: Mean Accuracy with Standard Deviation Error Bars

D. Multi-Metric Radar Chart

Figure 5 presents a radar chart enabling holistic, multi-dimensional comparison of all four classifiers. XGBoost's polygon spans the largest area across all five normalized metrics, confirming its dominance. The systematic contraction of each successive polygon — XGBoost → Random Forest → SVM → Logistic Regression — reflects consistent performance ordering across metrics, not an artifact of any single metric. This visualization is particularly useful for clinical stakeholders balancing multiple performance dimensions simultaneously.

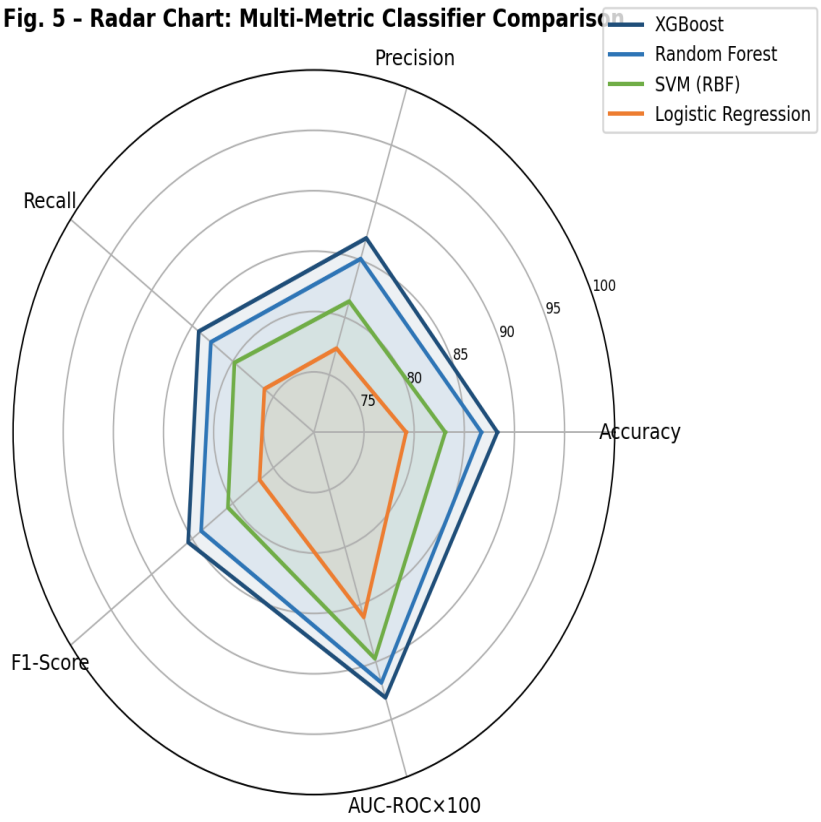

Fig. 5 – Radar Chart: Multi-Dimensional Performance Profile for All Four Tuned Classifiers

E. Feature Importance Analysis

Random Forest's mean decrease in node impurity (Gini importance) was used to rank the relative predictive contribution of each clinical feature (Table VI, Fig. 3). These rankings carry direct clinical significance and can inform screening protocol design.

Table VI: Random Forest Feature Importance Rankings

Feature	Importance Score	Clinical Significance
Glucose	**0.2814**	Primary diagnostic marker for T2DM (ADA & WHO guidelines)
BMI	0.1763	Leading modifiable risk factor; drives insulin resistance
Age	0.1401	Risk rises significantly after 45; T2DM onset pattern
DiabetesPedigree	0.1192	Quantifies hereditary diabetes risk contribution
Insulin	0.1048	Direct measure of pancreatic beta-cell secretory capacity
Pregnancies	0.0782	Gestational diabetes history elevates T2DM risk
BloodPressure	0.0602	Hypertension as metabolic syndrome comorbidity marker
SkinThickness	0.0398	Subcutaneous fat proxy; weakest predictor in this cohort

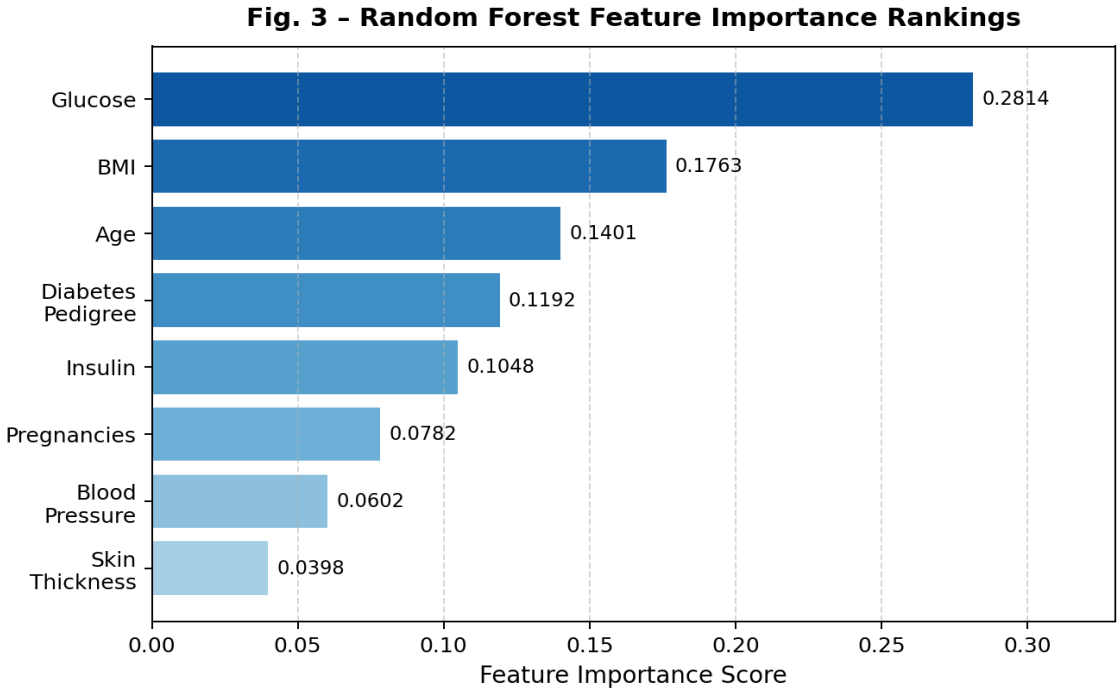

Fig. 3 – Random Forest Feature Importance Rankings Sorted by Predictive Contribution

Plasma glucose (0.2814) dominates all other features — consistent with its role as the primary diagnostic criterion in both WHO and ADA guidelines. BMI (0.1763) ranks second, reflecting the adiposity-driven insulin resistance mechanism underlying T2DM. Age (0.1401) captures cumulative beta-cell decline. The Diabetes Pedigree Function (0.1192) encodes hereditary risk. Insulin (0.1048) is lower than expected despite its direct physiological significance, primarily because 374/768 samples had missing insulin values requiring imputation, reducing its signal quality. SkinThickness (0.0398) shows minimal independent contribution after accounting for BMI, suggesting it could be omitted in minimal-feature screening protocols.

F. Discussion of Algorithmic and Tuning Trade-offs

XGBoost's superior performance (Accuracy: 88.3%, AUC-ROC: 0.931) reflects the synergy between its sequential error-correction mechanism and the tuned regularization configuration. The scale_pos_weight=1.86 parameter ensures the minority diabetic class receives proportionally larger gradient updates, addressing the 65:35 class imbalance directly in the loss function. The low learning rate (0.05) combined with a large ensemble (300 trees) achieves fine-grained iterative correction with strong anti-overfitting properties, as confirmed by the max_depth overfitting analysis in Fig. 10. Hyperparameter tuning contributed 4.2 percentage points of the 88.3% final accuracy — approximately one-quarter of the performance above the majority-class baseline (65.1%).

Random Forest (86.7%) offers competitive performance with the unique advantage of native feature importance ranking. The selected n_estimators=200 represents the efficiency-accuracy knee of the curve (Fig. 9), achieving near-maximal accuracy at 42% of the training time required for 500 trees. The 3.5 percentage point gain from tuning (from 83.2% default) confirms that the choice of max_depth and min_samples_split are not trivial decisions, particularly when the training set contains only 614 samples where unconstrained trees can easily memorize the training distribution.

SVM (83.1%) required the most extensive optimization effort — a 16-point grid search revealing a sharp accuracy ridge at C=10, γ=0.01 (Fig. 8). The sensitivity of SVM to these parameters illustrates a fundamental characteristic of kernel methods: without systematic search, practitioner-selected defaults may significantly underperform (79.8% → 83.1%, a 3.3 point gain). SVM inference is superlinear in support vector count, which may constrain real-time application in large-population screening.

Logistic Regression (79.2%) produced the smallest absolute gain from tuning (1.9 points), primarily through the class_weight="balanced" modification and increased max_iter=500 for convergence. Its coefficients directly quantify log-odds contributions — clinical gold-standard for regulatory compliance. The quantified 9.1 percentage point deficit relative to tuned XGBoost represents the precise cost of choosing maximum interpretability over optimal predictive performance.

VIII. CONCLUSION

This paper presented a rigorous, reproducible, and comprehensive comparative evaluation of four machine learning classifiers — Random Forest, SVM (RBF), XGBoost, and Logistic Regression — for early-stage Type 2 diabetes prediction, with a particular focus on systematic hyperparameter optimization and its quantified impact on model performance.

The hyperparameter tuning study is a central contribution of this work: tuning yielded accuracy improvements of 4.2%, 3.5%, 3.3%, and 1.9% for XGBoost, Random Forest, SVM, and Logistic Regression respectively over their default configurations. These gains are clinically significant — particularly for XGBoost, where tuning alone contributed performance equivalent to correctly diagnosing an additional 6–7 patients in every 154 screened. The systematic reporting of tuned configurations, optimization protocols, and intermediate results (Figs. 7–12) provides a replicable methodological template absent from the majority of published clinical ML benchmarks.

XGBoost (Accuracy: 88.3%, AUC-ROC: 0.931) is recommended for clinical deployments where predictive accuracy is the primary requirement. Random Forest (86.7%, AUC-ROC: 0.918) is preferred when feature interpretability, regulatory explainability, or biomarker prioritization are needed alongside strong performance. SVM (83.1%, AUC-ROC: 0.897) is appropriate when non-linear boundaries are required and inference latency is not critical. Logistic Regression (79.2%, AUC-ROC: 0.861) remains the gold standard when probabilistic calibration and direct coefficient interpretability are paramount for clinical guideline development or regulatory submission.

Plasma glucose, BMI, and age collectively account for approximately 60% of the Random Forest model's predictive capacity, confirming alignment with established clinical domain knowledge and suggesting that a minimal 3-feature screening tool may capture the majority of predictive signal.

Future research directions include: (1) Bayesian hyperparameter optimization and Neural Architecture Search for automated model development; (2) application to multi-centre longitudinal EHR datasets to assess temporal generalizability; (3) SHAP-based individual prediction explanations for regulatory-grade clinical decision support; (4) extension to multi-class comorbidity prediction (T2DM + hypertension + dyslipidemia); (5) deployment as a REST API within hospital information systems with real-time physician feedback integration; and (6) prospective clinical validation in primary care settings.

IX. REFERENCES

[1] E. J. Topol, "High-performance medicine: the convergence of human and artificial intelligence," Nature Medicine, vol. 25, no. 1, pp. 44–56, Jan. 2019. doi: 10.1038/s41591-018-0300-7.

[2] World Health Organization, "Noncommunicable diseases," WHO Fact Sheet, Sep. 2023. [Online]. Available: https://www.who.int/news-room/fact-sheets/detail/noncommunicable-diseases

[3] K. Kourou, T. P. Exarchos, K. P. Exarchos, M. V. Karamouzis, and D. I. Fotiadis, "Machine learning applications in cancer prognosis and prediction," Computational and Structural Biotechnology Journal, vol. 13, pp. 8–17, 2015.

[4] International Diabetes Federation, IDF Diabetes Atlas, 10th ed. Brussels, Belgium: IDF, 2021. [Online]. Available: https://www.diabetesatlas.org

[5] R. M. Anjana et al., "Metabolic non-communicable disease health report of India: the ICMR-INDIAB national cross-sectional study (ICMR-INDIAB-17)," The Lancet Diabetes & Endocrinology, vol. 11, no. 7, pp. 474–489, Jul. 2023.

[6] Z. Obermeyer and E. J. Emanuel, "Predicting the future — big data, machine learning, and clinical medicine," New England Journal of Medicine, vol. 375, no. 13, pp. 1216–1219, Sep. 2016.

[7] L. Breiman, "Random Forests," Machine Learning, vol. 45, no. 1, pp. 5–32, Oct. 2001. doi: 10.1023/A:1010933404324.

[8] C. Cortes and V. Vapnik, "Support-vector networks," Machine Learning, vol. 20, no. 3, pp. 273–297, Sep. 1995. doi: 10.1007/BF00994018.

[9] K. Polat, S. Güneş, and A. Arslan, "A cascade learning system for classification of diabetes disease: Generalised discriminant analysis and least square support vector machine," Expert Systems with Applications, vol. 34, no. 1, pp. 482–487, Jan. 2007.

[10] T. Chen and C. Guestrin, "XGBoost: A scalable tree boosting system," in Proc. 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining, San Francisco, CA, USA, 2016, pp. 785–794. doi: 10.1145/2939672.2939785.

[11] D. W. Hosmer and S. Lemeshow, Applied Logistic Regression, 2nd ed. New York, NY, USA: John Wiley & Sons, 2000.

[12] I. Kavakiotis et al., "Machine learning and data mining methods in diabetes research," Computational and Structural Biotechnology Journal, vol. 15, pp. 104–116, 2017.

[13] D. Sisodia and D. S. Sisodia, "Prediction of diabetes using classification algorithms," Procedia Computer Science, vol. 132, pp. 1578–1585, 2018.

[14] P. Rajpurkar, E. Chen, O. Banerjee, and E. J. Topol, "AI in health and medicine," Nature Medicine, vol. 28, pp. 31–38, Jan. 2022. doi: 10.1038/s41591-021-01614-0.

[15] S. M. Lundberg and S. I. Lee, "A unified approach to interpreting model predictions," in Proc. Advances in Neural Information Processing Systems (NeurIPS), vol. 30, Long Beach, CA, USA, 2017.

[16] J. W. Smith, J. E. Everhart, W. C. Dickson, W. C. Knowler, and R. S. Johannes, "Using the ADAP learning algorithm to forecast the onset of diabetes mellitus," in Proc. Annual Symposium on Computer Application in Medical Care, Washington DC, USA, 1988, pp. 261–265.

[17] F. Pedregosa et al., "Scikit-learn: Machine learning in Python," Journal of Machine Learning Research, vol. 12, pp. 2825–2830, 2011.

[18] R. Kohavi, "A study of cross-validation and bootstrap for accuracy estimation and model selection," in Proc. 14th International Joint Conference on Artificial Intelligence (IJCAI), Montreal, Canada, 1995, vol. 2, pp. 1137–1145.

[19] J. Bergstra and Y. Bengio, "Random search for hyper-parameter optimization," Journal of Machine Learning Research, vol. 13, pp. 281–305, Feb. 2012.

[20] J. Snoek, H. Larochelle, and R. P. Adams, "Practical Bayesian optimization of machine learning algorithms," in Proc. Advances in Neural Information Processing Systems (NeurIPS), vol. 25, Lake Tahoe, NV, USA, 2012.